



```
In [1]: import pandas as pd
data = pd.read_csv("C:/Users/bhava/OneDrive/Pictures/Desktop/GPTC/5th SEM/E
X = data.drop('Scores', axis=1)
y = data['Hours']
from sklearn.model_selection import train_test_split
X_train, X_temp, y_train, y_temp = train_test_split(X, y, test_size=0.3, ra
X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, test_size=0
print(X)
print(y)
```

	Hours
0	2.5
1	5.1
2	3.2
3	8.5
4	3.5
5	1.5
6	9.2
7	5.5
8	8.3
9	2.7
10	7.7
11	5.9
12	4.5
13	3.3
14	1.1
15	8.9
16	2.5
17	1.9
18	6.1
19	7.4
20	2.7
21	4.8
22	3.8
23	6.9
24	7.8
0	2.5
1	5.1
2	3.2
3	8.5
4	3.5
5	1.5
6	9.2
7	5.5
8	8.3
9	2.7
10	7.7
11	5.9
12	4.5
13	3.3
14	1.1
15	8.9
16	2.5
17	1.9
18	6.1
19	7.4
20	2.7
21	4.8
22	3.8
23	6.9
24	7.8

Name: Hours, dtype: float64

## Train set

```
In [2]: import pandas as pd
from sklearn.model_selection import train_test_split
data = pd.read_csv("C:/Users/bhava/OneDrive/Pictures/Desktop/GPTC/5th SEM/E
X = data.drop('Scores', axis=1)
y = data['Hours']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ra
X
```

Out[2]:

	Hours
0	2.5
1	5.1
2	3.2
3	8.5
4	3.5
5	1.5
6	9.2
7	5.5
8	8.3
9	2.7
10	7.7
11	5.9
12	4.5
13	3.3
14	1.1
15	8.9
16	2.5
17	1.9
18	6.1
19	7.4
20	2.7
21	4.8
22	3.8
23	6.9
24	7.8

## Validation sat


```
In [3]: import csv
from sklearn.model_selection import train_test_split
data = []
with open("C:/Users/bhava/OneDrive/Pictures/Desktop/GPTC/5th SEM/EXCEL/stud
csv_reader = csv.reader(file)
for row in csv_reader:
    data.append(row)
X = [row[:-1] for row in data]
y = [row[-1] for row in data]
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2, rand
X
```

```
Out[3]: [['Hours'],
['2.5'],
['5.1'],
['3.2'],
['8.5'],
['3.5'],
['1.5'],
['9.2'],
['5.5'],
['8.3'],
['2.7'],
['7.7'],
['5.9'],
['4.5'],
['3.3'],
['1.1'],
['8.9'],
['2.5'],
['1.9'],
['6.1'],
['7.4'],
['2.7'],
['4.8'],
['3.8'],
['6.9'],
['7.8']]
```

## Overfitting

```
In [4]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeRegressor
from sklearn.metrics import mean_squared_error
import matplotlib.pyplot as plt
data = pd.read_csv("C:/Users/bhava/OneDrive/Pictures/Desktop/GPTC/5th SEM/E
X = data.drop('Scores', axis=1)
y = data['Hours']
X.fillna(X.mean(), inplace=True)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ra
model = DecisionTreeRegressor(max_depth=None)
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
print(f"Mean Squared Error: {mse}")
plt.scatter(y_test, y_pred)
plt.xlabel("Actual Values")
plt.ylabel("Predicted Values")
plt.title("Actual vs. Predicted Values")
plt.show()
```


Mean Squared Error: 0.08199999999999992



## Underfitting

```
In [5]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
import matplotlib.pyplot as plt
data = pd.read_csv("C:/Users/bhava/OneDrive/Pictures/Desktop/GPTC/5th SEM/E
X = data.drop('Scores', axis=1)
y = data['Hours']
X.fillna(X.mean(), inplace=True)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ra
model = LinearRegression()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
print(f"Mean Squared Error: {mse}")
plt.scatter(y_test, y_pred)
plt.xlabel("Actual Values")
plt.ylabel("Predicted Values")
plt.title("Actual vs. Predicted Values")
plt.show()
```

Mean Squared Error: 1.262177448353619e-30




## Simple linear regression

```

In [6]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
import matplotlib.pyplot as plt
data = pd.read_csv("C:/Users/bhava/OneDrive/Pictures/Desktop/GPTC/5th SEM/E
X = data[['Number']]
y = data['Height']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ra
model = LinearRegression()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
plt.scatter(X, y, color='blue', label='Data')
plt.plot(X_test, y_pred, color='red', linewidth=2, label='Regression Line')
plt.xlabel('Independent Variable')
plt.ylabel('Dependent Variable')
plt.legend()
plt.show()
slope = model.coef_[0]
intercept = model.intercept_
print(f"Slope (m): {slope}")
print(f"Intercept (b): {intercept}")

```



Slope (m): -0.45626191596593185  
Intercept (b): 69.36866404251712

## Multiple linear regression¶



```
In [7]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
data = pd.read_csv("C:/Users/bhava/OneDrive/Pictures/Desktop/GPTC/5th SEM/E
print(data.head())
print(data.info())
X = data[['Scores']]
y = data['Hours']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ra
model = LinearRegression()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
print("Mean Squared Error:", mse)
print("R-squared:", r2)
print("Coefficients:", model.coef_)
print("Intercept:", model.intercept_)
```

```

      Hours  Scores
0      2.5      21
1      5.1      47
2      3.2      27
3      8.5      75
4      3.5      30
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 25 entries, 0 to 24
Data columns (total 2 columns):
 #   Column  Non-Null Count  Dtype
---  -
 0   Hours   25 non-null         float64
 1   Scores  25 non-null         int64
dtypes: float64(1), int64(1)
memory usage: 528.0 bytes
None
Mean Squared Error: 0.16754295656677842
R-squared: 0.9696127835594206
Coefficients: [0.09802864]
Intercept: -0.024756258820576527
```

## Polynomial linear regression

```


In [8]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import PolynomialFeatures
from sklearn.metrics import mean_squared_error, r2_score
import matplotlib.pyplot as plt

data = pd.read_csv("C:/Users/bhava/OneDrive/Pictures/Desktop/GPTC/5th SEM/E
X = data[['Hours']]
y = data['Scores']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ra
degree = 2 # Adjust the degree of the polynomial as needed
poly_features = PolynomialFeatures(degree=degree)
X_train_poly = poly_features.fit_transform(X_train)
X_test_poly = poly_features.transform(X_test)
poly_reg = LinearRegression()
poly_reg.fit(X_train_poly, y_train)
y_pred = poly_reg.predict(X_test_poly)
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
print("Mean Squared Error:", mse)
print("R-squared (R2) Score:", r2)
X_test_sorted, y_pred_sorted = zip(*sorted(zip(X_test.values, y_pred)))
plt.scatter(X_test, y_test, color='blue', label='Actual Data')
plt.plot(X_test_sorted, y_pred_sorted, color='red', label=f'Polynomial Regr
plt.xlabel('X')
plt.ylabel('y')
plt.legend()
plt.show()

```

Mean Squared Error: 21.066769638340965

R-squared (R2) Score: 0.9641965165901751



# logistic Regression¶

```
In [9]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
data = pd.read_csv("C:/Users/bhava/OneDrive/Pictures/Desktop/GPTC/5th SEM/E
X = data.drop('Hours', axis=1)
y = data['Scores']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ra
model = LogisticRegression()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy}")
plt.scatter(X_test, y_test, color='blue', label='Actual Data')
plt.xlabel('X')
plt.ylabel('y')
plt.legend()
plt.show()
```

Accuracy: 0.2

C:\Users\bhava\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.


Increase the number of iterations (max\_iter) or scale the data as shown i  
n:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression) ([https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression))

```
n_iter_i = _check_optimize_result(
```



## Equation regression model

```

In [10]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
import numpy as np
test_data = pd.read_csv("C:/Users/bhava/OneDrive/Pictures/Desktop/GPTC/5th
X_test = test_data[['Scores']]
y_test = test_data['Hours']
model = LinearRegression()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score

mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
mae = mean_absolute_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print("Mean Squared Error:", mse)
print("Root Mean Squared Error:", rmse)
print("Mean Absolute Error:", mae)
print("R-squared:", r2)

import matplotlib.pyplot as plt


plt.scatter(y_test, y_pred)
plt.xlabel("Actual Values")
plt.ylabel("Predicted Values")
plt.title("Actual vs. Predicted Values")
plt.show()

```

```

Mean Squared Error: 2659.5691999999995
Root Mean Squared Error: 51.57101123693426
Mean Absolute Error: 46.468
R-squared: -433.4951590052434

```



## Evaluation Matrix

```
In [11]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
import numpy as np
data = pd.read_csv("C:/Users/bhava/OneDrive/Pictures/Desktop/GPTC/5th SEM/E
X = data.drop('Hours', axis=1)
y = data['Scores']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, ra
model = LinearRegression()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
mae = mean_absolute_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
print(f'Mean Squared Error (MSE): {mse:.4f}')
print(f'Root Mean Squared Error (RMSE): {rmse:.4f}')
print(f'Mean Absolute Error (MAE): {mae:.4f}')
print(f'R-squared (R²): {r2:.4f}')
```

Mean Squared Error (MSE): 0.0000  
 Root Mean Squared Error (RMSE): 0.0000  
 Mean Absolute Error (MAE): 0.0000  
 R-squared (R²): 1.0000

## R square R2

```
In [12]: ##Root Mean Squared Error (RMSE)
import numpy as np
from sklearn.metrics import mean_squared_error
y_pred = model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
rmse
```

Out[12]: 0.0

## Route mean squared error

```
In [13]: ## adjusted R-squared
from sklearn.metrics import r2_score
r2 = r2_score(y_test, y_pred)
num_features = X_train.shape[1]
num_data_points = len(y_test)
adjusted_r2 = 1 - (1 - r2) * (num_data_points - 1) / (num_data_points - num
adjusted_r2
```

Out[13]: 1.0

## Mean absolute error

```
In [14]: ## Mean Absolute Error(MAE)
import pandas as pd
from sklearn.metrics import mean_absolute_error
data=pd.read_csv("C:/Users/bhava/OneDrive/Pictures/Desktop/GPTC/5th SEM/EXC
actual = data['Height']
predicted = data['Number']
mae = mean_absolute_error(actual, predicted)
print(f"Mean Absolute Error (MAE): {mae}")
```

Mean Absolute Error (MAE): 59.812004198866795

## Mean squared error

```
In [15]: ## Mean Squared Error
import pandas as pd
from sklearn.metrics import mean_squared_error
data = pd.read_csv("C:/Users/bhava/OneDrive/Pictures/Desktop/GPTC/5th SEM/E
actual = data['Number']
predicted = data['Height']
mse = mean_squared_error(actual, predicted)
print(f"Mean Squared Error (MSE): {mse}")
```

Mean Squared Error (MSE): 3592.3321704373852



# Cross validation

```
In [47]: import pandas as pd
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import KFold
from sklearn.ensemble import RandomForestRegressor # Example model
data = pd.read_csv("C:/Users/bhava/OneDrive/Pictures/Desktop/GPTC/5th SEM/E
X = data[['Height', 'Weight']] # Features
y = data['Number'] # Target variable
model = RandomForestRegressor() # Example model
kf = KFold(n_splits=5, shuffle=True, random_state=42)
mae_scores = -cross_val_score(model, X, y, cv=kf, scoring='neg_mean_absolut
mean_mae = mae_scores.mean()
std_mae = mae_scores.std()
```

```
In [48]: mae_scores
```

```
Out[48]: array([0.01570444, 0.00432278, 0.01376111, 0.01050611, 0.01677333])
```

```
In [49]: mean_mae
```

```
Out[49]: 0.012213555555822594
```

```
In [50]: std_mae
```

```
Out[50]: 0.004484474267726706
```

# Hold out method

```
In [18]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score
data=pd.read_csv("C:/Users/bhava/OneDrive/Pictures/Desktop/GPTC/5th SEM/EXC
X = data.drop('Hours', axis=1)
y = data['Scores']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, ra
model = DecisionTreeClassifier()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy: {accuracy}')
```

Accuracy: 0.125

# leave one out cross validation

```
In [20]: import pandas as pd
from sklearn.model_selection import LeaveOneOut
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score
data = pd.read_csv("C:/Users/bhava/OneDrive/Pictures/Desktop/GPTC/5th SEM/E
X = data.drop('Hours', axis=1)
y = data['Scores']
loo = LeaveOneOut()
predicted_labels = []
true_labels = []
for train_index, test_index in loo.split(X):
    X_train, X_test = X.iloc[train_index], X.iloc[test_index]
    y_train, y_test = y.iloc[train_index], y.iloc[test_index]
    model = DecisionTreeClassifier()
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)

    predicted_labels.extend(y_pred)
    true_labels.extend(y_test)
accuracy = accuracy_score(true_labels, predicted_labels)
print(f'Accuracy: {accuracy}')
```

Accuracy: 0.12

## k fold cross validation

```
In [21]: import pandas as pd
from sklearn.model_selection import cross_val_score, KFold
from sklearn.tree import DecisionTreeClassifier
data = pd.read_csv("C:/Users/bhava/OneDrive/Pictures/Desktop/GPTC/5th SEM/E
X = data.drop('Hours', axis=1)
y = data['Scores']
kfold = KFold(n_splits=5, shuffle=True, random_state=42)
model = DecisionTreeClassifier()
scores = cross_val_score(model, X, y, cv=kfold)

for i, score in enumerate(scores, start=1):
    print(f'Fold {i}: {score:.4f}')
print(f'Mean Accuracy: {scores.mean():.4f}')
print(f'Standard Deviation: {scores.std():.4f}')
```

Fold 1: 0.2000  
 Fold 2: 0.0000  
 Fold 3: 0.2000  
 Fold 4: 0.2000  
 Fold 5: 0.0000  
 Mean Accuracy: 0.1200  
 Standard Deviation: 0.0980

## Random forest regression

```
In [22]: from sklearn.ensemble import RandomForestRegressor
from sklearn.datasets import make_regression
X, y = make_regression(n_features=4, n_informative=2, random_state=0, shuffle=False)
rfr = RandomForestRegressor(max_depth=3)
rfr.fit(X, y)
print(rfr.predict([[0, 1, 0, 1]]))
```

```
[36.7820371]
```

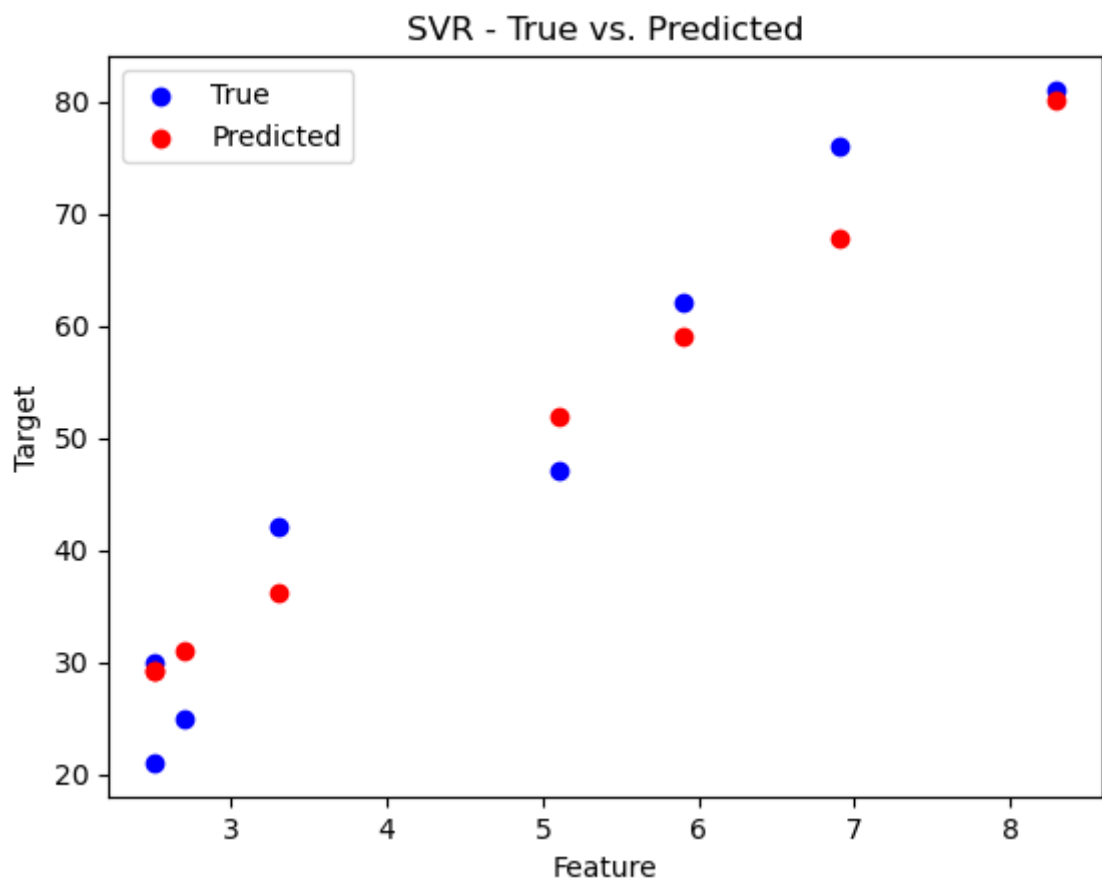
## Support vector regression

```

In [23]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.svm import SVR
from sklearn.metrics import mean_squared_error, r2_score
data = pd.read_csv("C:/Users/bhava/OneDrive/Pictures/Desktop/GPTC/5th SEM/E
X = data[['Hours']]
y = data['Scores']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, ra
model = SVR(kernel='linear')
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
plt.scatter(X_test, y_test, color='b', label='True')
plt.scatter(X_test, y_pred, color='r', label='Predicted')
plt.title('SVR - True vs. Predicted')
plt.xlabel('Feature')
plt.ylabel('Target')
plt.legend(loc='upper left')
plt.show()

print(f'Mean Squared Error (MSE): {mse:.4f}')
print(f'R-squared (R2): {r2:.4f}')

```



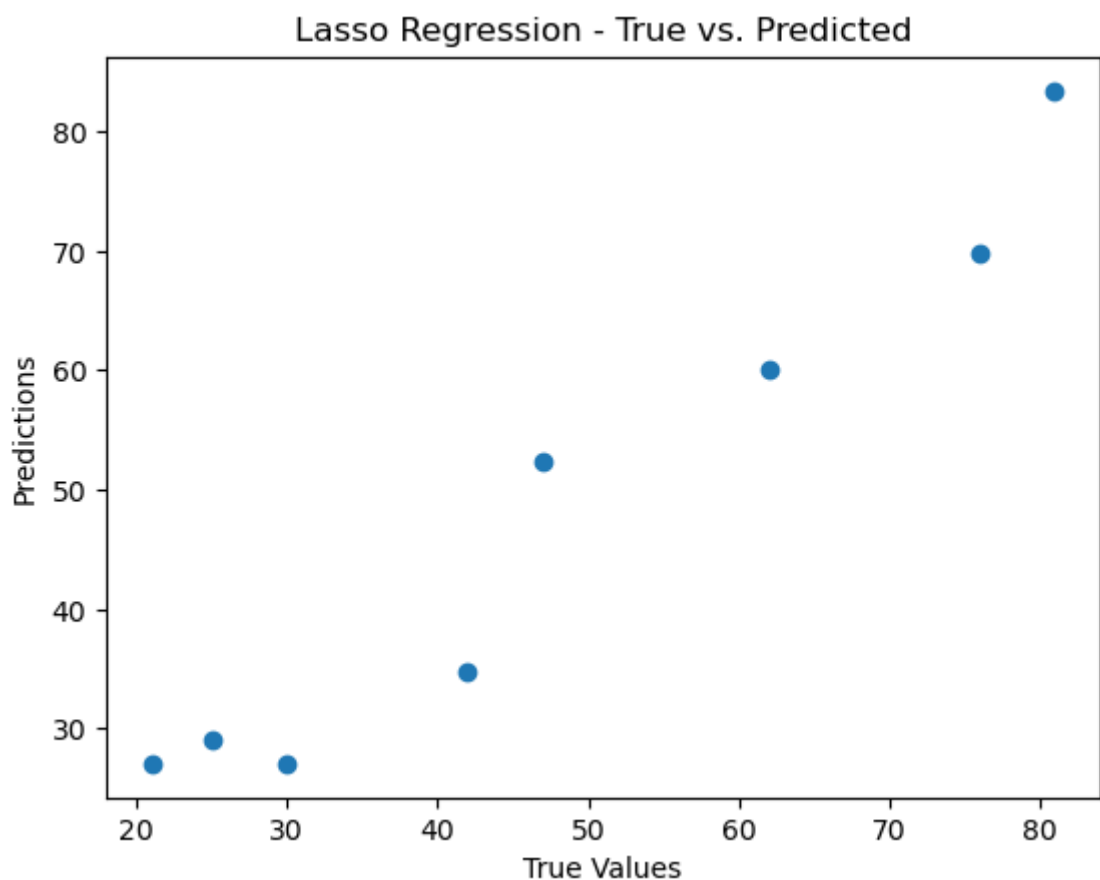
Mean Squared Error (MSE): 29.9038

R-squared (R<sup>2</sup>): 0.9351

# lasso Regression

```
In [24]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import Lasso
from sklearn.metrics import mean_squared_error, r2_score
data = pd.read_csv("C:/Users/bhava/OneDrive/Pictures/Desktop/GPTC/5th SEM/E
X = data[['Hours']]
y = data['Scores']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, ra
model = Lasso(alpha=0.01)
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
plt.scatter(y_test, y_pred)
plt.xlabel('True Values')
plt.ylabel('Predictions')
plt.title('Lasso Regression - True vs. Predicted')
plt.show()

print(f'Mean Squared Error (MSE): {mse:.4f}')
print(f'R-squared (R²): {r2:.4f}')
```



Mean Squared Error (MSE): 23.6231

R-squared ( $R^2$ ): 0.9488

# Decision tree classifier

```
In [25]: import pandas as pd
from sklearn.model_selection import LeaveOneOut
from sklearn.tree import DecisionTreeClassifier # Replace with your chosen
from sklearn.metrics import accuracy_score # Replace with appropriate eval
data = pd.read_csv("C:/Users/bhava/OneDrive/Pictures/Desktop/GPTC/5th SEM/E

X = data.drop('Hours', axis=1)
y = data['Scores']
loo = LeaveOneOut()
model = DecisionTreeClassifier()
predicted_labels = []
true_labels = []
for train_index, test_index in loo.split(X):
    X_train, X_test = X.iloc[train_index], X.iloc[test_index]
    y_train, y_test = y.iloc[train_index], y.iloc[test_index]

model.fit(X_train, y_train)
y_pred = model.predict(X_test)

predicted_labels.extend(y_pred)
true_labels.extend(y_test)
accuracy = accuracy_score(true_labels, predicted_labels)
print(f'Accuracy: {accuracy}')
```

Accuracy: 0.0

## Classification

```
In [26]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report
data = pd.read_csv("C:/Users/bhava/OneDrive/Pictures/Desktop/GPTC/5th SEM/E
X = data.drop('Hours', axis=1)
y = data['Scores']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ra
classifier = RandomForestClassifier(n_estimators=100, random_state=42)
classifier.fit(X_train, y_train)
y_pred = classifier.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.2f}")
print("Classification Report:")
print(classification_report(y_test, y_pred))
```

Accuracy: 0.20

Classification Report:

	precision	recall	f1-score	support
20	0.00	0.00	0.00	0
21	0.00	0.00	0.00	1
30	1.00	1.00	1.00	1
60	0.00	0.00	0.00	0
62	0.00	0.00	0.00	1
75	0.00	0.00	0.00	0
76	0.00	0.00	0.00	1
81	0.00	0.00	0.00	1
85	0.00	0.00	0.00	0
accuracy			0.20	5
macro avg	0.11	0.11	0.11	5
weighted avg	0.20	0.20	0.20	5

```

C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))

```

## Binary classifier



```
In [27]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report
data = pd.read_csv("C:/Users/bhava/OneDrive/Pictures/Desktop/GPTC/5th SEM/E
X = data.iloc[:, :-1]
y = data.iloc[:, -1]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ra
logistic_classifier = LogisticRegression()
logistic_classifier.fit(X_train, y_train)
y_pred = logistic_classifier.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.2f}")
print("Classification Report:")
print(classification_report(y_test, y_pred))
```

Accuracy: 0.20

Classification Report:

	precision	recall	f1-score	support
21	0.00	0.00	0.00	1
30	0.50	1.00	0.67	1
62	0.00	0.00	0.00	1
67	0.00	0.00	0.00	0
69	0.00	0.00	0.00	0
76	0.00	0.00	0.00	1
81	0.00	0.00	0.00	1
88	0.00	0.00	0.00	0
accuracy			0.20	5
macro avg	0.06	0.12	0.08	5
weighted avg	0.10	0.20	0.13	5

```
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression) ([https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression))

```
n_iter_i = _check_optimize_result(
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
```

## Multi class classifier

```
In [28]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, classification_report
data = pd.read_csv("C:/Users/bhava/OneDrive/Pictures/Desktop/GPTC/5th SEM/E
X = data.iloc[:, :-1]
y = data.iloc[:, -1]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ra
svm_classifier = SVC(kernel='linear')
svm_classifier.fit(X_train, y_train)
y_pred = svm_classifier.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.2f}")
print("Classification Report:")
print(classification_report(y_test, y_pred))
```

Accuracy: 0.20

Classification Report:

	precision	recall	f1-score	support
21	0.00	0.00	0.00	1
30	0.50	1.00	0.67	1
62	0.00	0.00	0.00	1
67	0.00	0.00	0.00	0
69	0.00	0.00	0.00	0
75	0.00	0.00	0.00	0
76	0.00	0.00	0.00	1
81	0.00	0.00	0.00	1
accuracy			0.20	5
macro avg	0.06	0.12	0.08	5
weighted avg	0.10	0.20	0.13	5

```

C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))

```

## Confusion matrix

```
In [29]: import pandas as pd
from sklearn import model_selection
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix
data = pd.read_csv("C:/Users/bhava/OneDrive/Pictures/Desktop/GPTC/5th SEM/E
# Split the data into features (X) and the target variable (y)
X = data.drop('Hours', axis=1) # Adjust 'target_column' to your target var
y = data['Scores']
test_size = 0.33
X_train, X_test, Y_train, Y_test = model_selection.train_test_split(X, y, t
model = LogisticRegression(solver='liblinear')
model.fit(X_train, Y_train)
predicted = model.predict(X_test)
matrix = confusion_matrix(Y_test, predicted)
print(matrix)
```

```
[[0 0 0 1 0 0 0 0 0 0]
 [0 0 0 1 0 0 0 0 0 0]
 [0 0 0 1 0 0 0 0 0 0]
 [0 0 0 1 0 0 0 0 0 0]
 [0 0 0 0 0 0 0 0 0 1]
 [0 0 0 0 0 0 0 0 0 1]
 [0 0 0 0 0 0 0 0 0 1]
 [0 0 0 0 0 0 0 0 0 1]
 [0 0 0 0 0 0 0 0 0 1]
 [0 0 0 0 0 0 0 0 0 1]
 [0 0 0 0 0 0 0 0 0 0]]
```

## AUC ROC

```
In [31]: import pandas as pd
data = pd.read_csv("C:/Users/bhava/OneDrive/Pictures/Desktop/GPTC/5th SEM/E
column_name = 'Scores'

# Calculate Gini index for a specific column
def calculate_gini(data, column):
    total_rows = len(data)
    unique_classes = data[column].unique()
    gini = 0

    for label in unique_classes:
        proportion = (len(data[data[column] == label]) / total_rows) ** 2
        gini += proportion

    gini_index = 1 - gini
    return gini_index

gini_value = calculate_gini(data, column_name)
print(f"Gini Index for '{column_name}': {gini_value:.4f}")
```

Gini Index for 'Scores': 0.9504

## Gini

```
In [ ]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import roc_auc_score
data = pd.read_csv("C:/Users/bhava/OneDrive/Pictures/Desktop/GPTC/5th SEM/E
X = data.iloc[:, :-1]
y = data.iloc[:, -1]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ra
rf_classifier = RandomForestClassifier(n_estimators=100, random_state=42)
rf_classifier.fit(X_train, y_train)
y_prob = rf_classifier.predict_proba(X_test)[:, 1] # Considering positive
auc_roc = roc_auc_score(y_test, y_prob)
print(f"AUC-ROC Score: {auc_roc:.2f}")
```

## Entropy

```
In [32]: import pandas as pd
import math
data = pd.read_csv("C:/Users/bhava/OneDrive/Pictures/Desktop/GPTC/5th SEM/E
column_name = 'Hours'

# Calculate entropy for a specific column
def calculate_entropy(data, column):
    total_rows = len(data)
    unique_classes = data[column].unique()
    entropy = 0

    for label in unique_classes:
        proportion = len(data[data[column] == label]) / total_rows
        entropy -= proportion * math.log2(proportion)

    return entropy

entropy_value = calculate_entropy(data, column_name)
print(f"Entropy for '{column_name}': {entropy_value:.4f}")
```

Entropy for 'Hours': 4.4839

## Boston housing price from sci-kit learn

```
In [33]: from sklearn.datasets import load_boston
boston = load_boston()
X = boston.data
y = boston.target
print(boston.DESCR)
```

```
.. _boston_dataset:
```

```
Boston house prices dataset
```

```
-----
```

```
**Data Set Characteristics:**
```

```
    :Number of Instances: 506
```

```
    :Number of Attributes: 13 numeric/categorical predictive. Median Value  
    (attribute 14) is usually the target.
```

```
    :Attribute Information (in order):
```

```
        - CRIM      per capita crime rate by town  
        - ZN        proportion of residential land zoned for lots over 25,0  
00 sq.ft.  
        - INDUS     proportion of non-retail business acres per town  
        - CHAS      Charles River dummy variable (= 1 if tract bounds rive  
r; 0 otherwise)  
        - NOX       nitric oxides concentration (parts per 10 million)  
        - RM        average number of rooms per dwelling  
        - AGE       proportion of owner-occupied units built prior to 1940  
        - DIS       weighted distances to five Boston employment centres  
        - RAD       index of accessibility to radial highways  
        - TAX       full-value property-tax rate per $10,000  
        - PTRATIO   pupil-teacher ratio by town  
        - B          $1000(B_k - 0.63)^2$  where  $B_k$  is the proportion of black p  
eople by town  
        - LSTAT     % lower status of the population  
        - MEDV      Median value of owner-occupied homes in $1000's
```

```
    :Missing Attribute Values: None
```

```
    :Creator: Harrison, D. and Rubinfeld, D.L.
```

```
This is a copy of UCI ML housing dataset.
```

```
https://archive.ics.uci.edu/ml/machine-learning-databases/housing/ (http  
s://archive.ics.uci.edu/ml/machine-learning-databases/housing/)
```

```
This dataset was taken from the StatLib library which is maintained at Car  
negie Mellon University.
```

```
The Boston house-price data of Harrison, D. and Rubinfeld, D.L. 'Hedonic  
prices and the demand for clean air', J. Environ. Economics & Management,  
vol.5, 81-102, 1978. Used in Belsley, Kuh & Welsch, 'Regression diagnost  
ics
```

```
...', Wiley, 1980. N.B. Various transformations are used in the table on  
pages 244-261 of the latter.
```

```
The Boston house-price data has been used in many machine learning papers  
that address regression  
problems.
```

```
.. topic:: References
```

```
    - Belsley, Kuh & Welsch, 'Regression diagnostics: Identifying Influenti  
al Data and Sources of Collinearity', Wiley, 1980. 244-261.
```

```
    - Quinlan,R. (1993). Combining Instance-Based and Model-Based Learning.  
    In Proceedings on the Tenth International Conference of Machine Learning,
```



```
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\utils\deprecation.py:8
7: FutureWarning: Function load_boston is deprecated; `load_boston` is dep
recated in 1.0 and will be removed in 1.2.
```

The Boston housing prices dataset has an ethical problem. You can refer to the documentation of this function for further details.

The scikit-learn maintainers therefore strongly discourage the use of this dataset unless the purpose of the code is to study and educate about ethical issues in data science and machine learning.

In this special case, you can fetch the dataset from the original source::

```
import pandas as pd
import numpy as np

data_url = "http://lib.stat.cmu.edu/datasets/boston"
raw_df = pd.read_csv(data_url, sep="\s+", skiprows=22, header=None)
e) data = np.hstack([raw_df.values[::2, :], raw_df.values[1::2, :2]])
target = raw_df.values[1::2, 2]
```

Alternative datasets include the California housing dataset (i.e. :func:`~sklearn.datasets.fetch\_california\_housing`) and the Ames housing dataset. You can load the datasets as follows::

```
from sklearn.datasets import fetch_california_housing
housing = fetch_california_housing()

for the California housing dataset and::

from sklearn.datasets import fetch_openml
housing = fetch_openml(name="house_prices", as_frame=True)

for the Ames housing dataset.
```

```
warnings.warn(msg, category=FutureWarning)
```

```
In [34]: from sklearn.datasets import load_boston
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
boston = load_boston()
X = boston.data
y = boston.target
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ra
model = LinearRegression()
model.fit(X_train, y_train)
predictions = model.predict(X_test)
mse = mean_squared_error(y_test, predictions)
r2 = r2_score(y_test, predictions)

print("Mean Squared Error (MSE):", mse)
print("R-squared (R2):", r2)
```

Mean Squared Error (MSE): 24.29111947497374  
R-squared (R2): 0.6687594935356289

## Cricket match result - past data

```
In [35]: import pandas as pd

data = pd.read_csv("C:/Users/bhava/OneDrive/Pictures/Desktop/GPTC/5th SEM/E
print(data.head())

run_a_wins = len(data[data['Runs'] == 'run A'])
hs_b_wins = len(data[data['HS'] == 'highest_scores B'])

print("Number of matches won by Team A:", run_a_wins)
print("Number of matches won by Team B:", hs_b_wins)
```

	Unnamed: 0		Player	Span	Mat	Inns	NO	Runs
\								
0	0		SR Tendulkar (INDIA)	1989-2012	463	452	41	18426
1	1		KC Sangakkara (Asia/ICC/SL)	2000-2015	404	380	41	14234
2	2		RT Ponting (AUS/ICC)	1995-2012	375	365	39	13704
3	3		ST Jayasuriya (Asia/SL)	1989-2011	445	433	18	13430
4	4		DPMD Jayawardene (Asia/SL)	1998-2015	448	418	39	12650

	HS	Ave	BF	SR	100	50	0	Unnamed: 13
0	200*	44.83	21367	86.23	49	96	20	NaN
1	169	41.98	18048	78.86	25	93	15	NaN
2	164	42.03	17046	80.39	30	82	20	NaN
3	189	32.36	14725	91.2	28	68	34	NaN
4	144	33.37	16020	78.96	19	77	28	NaN

Number of matches won by Team A: 0  
Number of matches won by Team B: 0

## Cricket match result-past data

```
In [36]: import pandas as pd
from sklearn.linear_model import LinearRegression
import matplotlib.pyplot as plt
data = pd.read_csv('integrated_data.csv')
print(data.head())
X = data[['average_rain_fall_mm_per_year_y']]
y = data['hg/ha_yield']
model = LinearRegression()
model.fit(X, y)
predicted_yield = model.predict(X)
plt.scatter(X, y, color='blue', label='Actual data')
plt.plot(X, predicted_yield, color='red', label='Predicted yield')
plt.xlabel('Rainfall')
plt.ylabel('Yield')
plt.legend()
plt.title('Crop Yield vs Rainfall')
plt.show()
```


C:\Users\bhava\AppData\Local\Temp\ipykernel\_13236\3090302346.py:4: DtypeWarning: Columns (2) have mixed types. Specify dtype option on import or set low\_memory=False.

```
data = pd.read_csv('integrated_data.csv')
```

	Area	Year	average_rain_fall_mm_per_year_x	Unnamed: 0	Area
0	Afghanistan	1990	327	0	Albania
1	Afghanistan	1990	327	1	Albania
2	Afghanistan	1990	327	2	Albania
3	Afghanistan	1990	327	3	Albania
4	Afghanistan	1990	327	4	Albania

	Item	hg/ha_yield	average_rain_fall_mm_per_year_y	\
0	Maize	36613	1485.0	
1	Potatoes	66667	1485.0	
2	Rice, paddy	23333	1485.0	
3	Sorghum	12500	1485.0	
4	Soybeans	7000	1485.0	

	pesticides_tonnes	avg_temp
0	121.0	16.37
1	121.0	16.37
2	121.0	16.37
3	121.0	16.37
4	121.0	16.37



**Crope yield pastdata**

```
In [37]: from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
import pandas as pd
iris = load_iris()
iris_df = pd.DataFrame(data=iris.data, columns=iris.feature_names)
iris_df['target'] = iris.target
print(iris_df.head())
X = iris.data # Features
y = iris.target # Target variable
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ra
print("X_train shape:", X_train.shape)
print("X_test shape:", X_test.shape)
print("y_train shape:", y_train.shape)
print("y_test shape:", y_test.shape)
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)
0	5.1	3.5	1.4	0.
1	4.9	3.0	1.4	0.
2	4.7	3.2	1.3	0.
3	4.6	3.1	1.5	0.
4	5.0	3.6	1.4	0.

	target
0	0
1	0
2	0
3	0
4	0

X\_train shape: (120, 4)  
X\_test shape: (30, 4)  
y\_train shape: (120,)  
y\_test shape: (30,)

## Iris dataset from sci-kit learn perform data exploration preprocessing and splitting

```
In [38]: from sklearn.datasets import load_breast_cancer
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, classification_report
breast_cancer = load_breast_cancer()
X = breast_cancer.data # Features
y = breast_cancer.target # Target variable
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ra
clf = DecisionTreeClassifier(random_state=42)
clf.fit(X_train, y_train)
y_pred = clf.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.2f}")
print("Classification Report:")
print(classification_report(y_test, y_pred))
```

Accuracy: 0.95

Classification Report:

	precision	recall	f1-score	support
0	0.93	0.93	0.93	43
1	0.96	0.96	0.96	71
accuracy			0.95	114
macro avg	0.94	0.94	0.94	114
weighted avg	0.95	0.95	0.95	114

## Build random forest based model

```
In [39]: from sklearn.datasets import load_breast_cancer
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report
breast_cancer = load_breast_cancer()
X = breast_cancer.data # Features
y = breast_cancer.target # Target variable
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ra
rf_classifier = RandomForestClassifier(n_estimators=100, random_state=42)
rf_classifier.fit(X_train, y_train)
y_pred = rf_classifier.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.2f}")
print("Classification Report:")
print(classification_report(y_test, y_pred))
```

Accuracy: 0.96

Classification Report:

	precision	recall	f1-score	support
0	0.98	0.93	0.95	43
1	0.96	0.99	0.97	71
accuracy			0.96	114
macro avg	0.97	0.96	0.96	114
weighted avg	0.97	0.96	0.96	114

## Max Voting

```

In [40]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.metrics import accuracy_score, classification_report
data = pd.read_csv("C:/Users/bhava/OneDrive/Pictures/Desktop/GPTC/5th SEM/E
X = data.iloc[:, :-1]
y = data.iloc[:, -1]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ra
rf_classifier = RandomForestClassifier(n_estimators=100, random_state=42)
gb_classifier = GradientBoostingClassifier(n_estimators=100, learning_rate=

rf_classifier.fit(X_train, y_train)
gb_classifier.fit(X_train, y_train)
y_pred_rf = rf_classifier.predict(X_test)
y_pred_gb = gb_classifier.predict(X_test)
voting_predictions = []

for i in range(len(X_test)):
    combined_prediction = max(y_pred_rf[i], y_pred_gb[i])
    voting_predictions.append(combined_prediction)
accuracy = accuracy_score(y_test, voting_predictions)
print(f"Accuracy: {accuracy:.2f}")
print("Classification Report:")
print(classification_report(y_test, voting_predictions))

```

Accuracy: 0.20

Classification Report:

	precision	recall	f1-score	support
21	0.00	0.00	0.00	1
30	0.50	1.00	0.67	1
62	0.00	0.00	0.00	1
67	0.00	0.00	0.00	0
69	0.00	0.00	0.00	0
75	0.00	0.00	0.00	0
76	0.00	0.00	0.00	1
81	0.00	0.00	0.00	1
accuracy			0.20	5
macro avg	0.06	0.12	0.08	5
weighted avg	0.10	0.20	0.13	5



```

C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))

```

## Averaging

```

In [41]: import numpy as np
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ra
from sklearn.metrics import accuracy_score, classification_report
data = pd.read_csv("C:/Users/bhava/OneDrive/Pictures/Desktop/GPTC/5th SEM/E
X = data.iloc[:, :-1]
y = data.iloc[:, -1]
predictions_model1 = np.array([0, 1, 1, 0, 1]) # Replace with actual predi
predictions_model2 = np.array([1, 1, 0, 0, 1]) # Replace with actual predi
predictions_model3 = np.array([1, 0, 1, 1, 0]) # Replace with actual predi

# Combine predictions using averaging
averaged_predictions = (predictions_model1 + predictions_model2 + predictio

# Round the averaged predictions for binary classification
rounded_averaged_predictions = np.round(averaged_predictions).astype(int)

# Display the averaged predictions and the rounded predictions
print("Averaged Predictions:", averaged_predictions)
print("Rounded Averaged Predictions:", rounded_averaged_predictions)
accuracy = accuracy_score(y_test, voting_predictions)
print(f"Accuracy: {accuracy:.2f}")

```

Averaged Predictions: [0.66666667 0.66666667 0.66666667 0.33333333 0.66666667]

Rounded Averaged Predictions: [1 1 1 0 1]

Accuracy: 0.20

## Weighted Average

```

In [42]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.metrics import accuracy_score, classification_report
data = pd.read_csv("C:/Users/bhava/OneDrive/Pictures/Desktop/GPTC/5th SEM/E
X = data.iloc[:, :-1]
y = data.iloc[:, -1]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ra
rf_classifier = RandomForestClassifier(n_estimators=100, random_state=42)
gb_classifier = GradientBoostingClassifier(n_estimators=100, learning_rate=

rf_classifier.fit(X_train, y_train)
gb_classifier.fit(X_train, y_train)
y_pred_rf = rf_classifier.predict(X_test)
y_pred_gb = gb_classifier.predict(X_test)
weight_rf = 0.6 # Weight for Random Forest's prediction
weight_gb = 0.4 # Weight for Gradient Boosting's prediction
weighted_average_predictions = (weight_rf * y_pred_rf) + (weight_gb * y_pre

accuracy = accuracy_score(y_test, weighted_average_predictions)
print(f"Accuracy: {accuracy:.2f}")

print("Classification Report:")
print(classification_report(y_test, weighted_average_predictions))

```

Accuracy: 0.00

Classification Report:

	precision	recall	f1-score	support
21.0	0.00	0.00	0.00	1.0
28.0	0.00	0.00	0.00	0.0
30.0	0.00	0.00	0.00	1.0
62.0	0.00	0.00	0.00	1.0
67.0	0.00	0.00	0.00	0.0
69.0	0.00	0.00	0.00	0.0
75.0	0.00	0.00	0.00	0.0
76.0	0.00	0.00	0.00	1.0
81.0	0.00	0.00	0.00	1.0
accuracy			0.00	5.0
macro avg	0.00	0.00	0.00	5.0
weighted avg	0.00	0.00	0.00	5.0

```

C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))

```

## Stacking

```

In [43]: from sklearn.ensemble import StackingClassifier
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
X, y = make_classification(n_samples=1000, n_features=20, random_state=42)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
base_estimators = [
    ('rf', RandomForestClassifier(n_estimators=100, random_state=42)),
    ('gb', GradientBoostingClassifier(n_estimators=100, learning_rate=0.1,
                                     random_state=42)),
]
meta_estimator = LogisticRegression()
stacking_classifier = StackingClassifier(estimators=base_estimators, final_estimator=meta_estimator)
stacking_classifier.fit(X_train, y_train)
predictions = stacking_classifier.predict(X_test)
accuracy = accuracy_score(y_test, predictions)
print(f"Accuracy of the stacked model: {accuracy:.2f}")

```

Accuracy of the stacked model: 0.90

## Blending

```

In [44]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.metrics import accuracy_score, classification_report
data = pd.read_csv("C:/Users/bhava/OneDrive/Pictures/Desktop/GPTC/5th SEM/E
X = data.iloc[:, :-1]
y = data.iloc[:, -1]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ra

rf_classifier = RandomForestClassifier(n_estimators=100, random_state=42)
gb_classifier = GradientBoostingClassifier(n_estimators=100, learning_rate=

rf_classifier.fit(X_train, y_train)
gb_classifier.fit(X_train, y_train)

y_pred_rf = rf_classifier.predict(X_test)
y_pred_gb = gb_classifier.predict(X_test)

blended_predictions = (y_pred_rf + y_pred_gb) / 2

final_prediction = [1 if pred >= 0.5 else 0 for pred in blended_predictions

accuracy = accuracy_score(y_test, final_prediction)
print(f"Accuracy: {accuracy:.2f}")

print("Classification Report:")
print(classification_report(y_test, final_prediction))

```

Accuracy: 0.00

Classification Report:

	precision	recall	f1-score	support
1	0.00	0.00	0.00	0.0
21	0.00	0.00	0.00	1.0
30	0.00	0.00	0.00	1.0
62	0.00	0.00	0.00	1.0
76	0.00	0.00	0.00	1.0
81	0.00	0.00	0.00	1.0
accuracy			0.00	5.0
macro avg	0.00	0.00	0.00	5.0
weighted avg	0.00	0.00	0.00	5.0

```

C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))

```

## Bagging

```
In [45]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report
data = pd.read_csv("C:/Users/bhava/OneDrive/Pictures/Desktop/GPTC/5th SEM/E
X = data.iloc[:, :-1]
y = data.iloc[:, -1]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ra

rf_classifier = RandomForestClassifier(n_estimators=100, random_state=42)

rf_classifier.fit(X_train, y_train)

y_pred = rf_classifier.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.2f}")

print("Classification Report:")
print(classification_report(y_test, y_pred))
```

Accuracy: 0.20

Classification Report:

	precision	recall	f1-score	support
21	0.00	0.00	0.00	1
30	0.50	1.00	0.67	1
62	0.00	0.00	0.00	1
67	0.00	0.00	0.00	0
69	0.00	0.00	0.00	0
75	0.00	0.00	0.00	0
76	0.00	0.00	0.00	1
81	0.00	0.00	0.00	1
accuracy			0.20	5
macro avg	0.06	0.12	0.08	5
weighted avg	0.10	0.20	0.13	5

```

C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))

```

## Boosting



```
In [46]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.metrics import accuracy_score, classification_report
data = pd.read_csv("C:/Users/bhava/OneDrive/Pictures/Desktop/GPTC/5th SEM/E
X = data.iloc[:, :-1]
y = data.iloc[:, -1]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ra
gb_classifier = GradientBoostingClassifier(n_estimators=100, learning_rate=

gb_classifier.fit(X_train, y_train)

y_pred = gb_classifier.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.2f}")

print("Classification Report:")
print(classification_report(y_test, y_pred))
```

Accuracy: 0.00

Classification Report:

	precision	recall	f1-score	support
21	0.00	0.00	0.00	1.0
25	0.00	0.00	0.00	0.0
30	0.00	0.00	0.00	1.0
62	0.00	0.00	0.00	1.0
67	0.00	0.00	0.00	0.0
69	0.00	0.00	0.00	0.0
75	0.00	0.00	0.00	0.0
76	0.00	0.00	0.00	1.0
81	0.00	0.00	0.00	1.0
accuracy			0.00	5.0
macro avg	0.00	0.00	0.00	5.0
weighted avg	0.00	0.00	0.00	5.0

```
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\bhava\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
```

In [ ]:

In [ ]: